



deeplearning.ai

Optimization
Algorithms

Mini-batch
gradient descent



Batch vs. mini-batch gradient descent

Vectorization allows you to efficiently compute on m examples.

$$X = \begin{bmatrix} x^{(1)} & x^{(2)} & x^{(3)} & \dots & x^{(1000)} & | & x^{(1001)} & \dots & x^{(2000)} & | & \dots & | & \dots & x^{(m)} \end{bmatrix}$$

(n_x, m) $X^{\{1\}} (n_x, 1000)$ $X^{\{2\}} (n_x, 1000)$ $X^{\{5,000\}} (n_x, 1000)$

$$Y = \begin{bmatrix} y^{(1)} & y^{(2)} & y^{(3)} & \dots & y^{(1000)} & | & y^{(1001)} & \dots & y^{(2000)} & | & \dots & | & \dots & y^{(m)} \end{bmatrix}$$

$(1, m)$ $Y^{\{1\}} (1, 1000)$ $Y^{\{2\}} (1, 1000)$ $Y^{\{5,000\}} (1, 1000)$



What if $m = 5,000,000$?

5,000 mini-batches of 1,000 each

Mini-batch t : $X^{\{t\}}, Y^{\{t\}}$

$$\begin{bmatrix} x^{(i)} \\ z^{[L]} \\ X^{\{t\}}, Y^{\{t\}} \end{bmatrix}$$

Mini-batch gradient descent

Repeat $\{$
 for $t = 1, \dots, 5000 \}$

Forward prop on X^{test} .

$$z^{(i)} = W^{(i)} X^{\text{test}} + b^{(i)}$$

$$A^{(i)} = g^{(i)}(z^{(i)})$$

$$\vdots$$

$$A^{(L)} = g^{(L)}(z^{(L)})$$

Vectorized implementation
(1000 examples)

for $X^{\text{test}}, Y^{\text{test}}$.

Compute cost $J^{\text{test}} = \frac{1}{1000} \sum_{i=1}^L \ell(\hat{y}^{(i)}, y^{(i)}) + \frac{\lambda}{2 \cdot 1000} \sum_{i=1}^L \|W^{(i)}\|_F^2$.

Backprop to compute gradients w.r.t J^{test} (using $(X^{\text{test}}, Y^{\text{test}})$)

$$W := W^{(i)} - \alpha dW^{(i)}, \quad b := b^{(i)} - \alpha db^{(i)}$$

"1 epoch"

pass through training set.

1 step of gradient descent
 using $X^{\text{test}}, Y^{\text{test}}$.
 (or if $m=1000$)

X, Y



deeplearning.ai

Optimization Algorithms


Understanding mini-batch gradient descent




Training with mini batch gradient descent



Batch gradient descent



Mini-batch gradient descent




Choosing your mini-batch size

☞ → If mini-batch size = m : Batch gradient descent. $(X^{(13)}, Y^{(13)}) = (X, Y)$.

☞ → If mini-batch size = 1 : Stochastic gradient descent. Every example is its own mini-batch.
 $(X^{(13)}, Y^{(13)}) = (x^{(1)}, y^{(1)}) \dots (x^{(n)}, y^{(n)})$ mini-batch.

In practice: Somewhere in-between 1 and m



Stochastic
gradient
descent



Less speedup
from vectorization

In-between
(mini-batch size
not too big/small)



Fastest learning.

- Vectorization.
(~ 1000)
- Make passes without
processing entire training set.

Batch
gradient descent
(mini-batch size = m)



Too long
per iteration

Choosing your mini-batch size

☐ If small toy set : Use batch gradient descent.
($m \leq 2000$)

☐ Typical mini-batch sizes:

→ 64, 128, 256, 512, 1024
 $2^6, 2^7, 2^8, 2^9, 2^{10}$

Make sure mini-batch fits in CPU/GPU memory.
 $X^{(t)}, Y^{(t)}$



deeplearning.ai

Optimization
Algorithms

Exponentially
weighted averages



Temperature in London

$$\theta_1 = 40^\circ\text{F} \quad 4^\circ\text{C} \leftarrow$$

$$\theta_2 = 49^\circ\text{F} \quad 9^\circ\text{C}$$


$$\theta_3 = 45^\circ\text{F} \quad \vdots$$

\vdots

$$\theta_{180} = 60^\circ\text{F} \quad 15^\circ\text{C}$$

$$\theta_{181} = 56^\circ\text{F} \quad \vdots$$

\vdots



$$V_0 = 0$$

$$V_1 = 0.9 V_0 + 0.1 \theta_1$$

$$V_2 = 0.9 V_1 + 0.1 \theta_2$$

$$V_3 = 0.9 V_2 + 0.1 \theta_3$$

\vdots

$$V_t = 0.9 V_{t-1} + 0.1 \theta_t$$



2:31 / 5:58



Andrew Ng


Exponentially weighted ^{moving} averages

☐
$$V_t = \beta V_{t-1} + (1-\beta) \theta_t \leftarrow$$

 $\beta = 0.9 : \approx 10 \text{ days' temper.}$
 $\beta = 0.98 : \approx 50 \text{ days}$
 $\beta = 0.5 : \approx 2 \text{ days}$

V_t is approximately
average over
 $\rightarrow \approx \frac{1}{1-\beta}$ days'
temperature.

$$\frac{1}{1-0.98} = 50$$





deeplearning.ai

Optimization Algorithms

Understanding exponentially weighted averages



Exponentially weighted averages

☞ $v_t = \beta v_{t-1} + (1 - \beta)\theta_t$

$$v_{100} = 0.9v_{99} + 0.1\theta_{100}$$

$$v_{99} = 0.9v_{98} + 0.1\theta_{99}$$

$$v_{98} = 0.9v_{97} + 0.1\theta_{97}$$

...


$$\begin{aligned} \rightarrow v_{100} &= 0.1\theta_{100} + 0.9 \cancel{v_{99}} (0.1\theta_{99} + 0.9 \cancel{v_{98}}) \\ &= \underbrace{0.1\theta_{100}} + \underbrace{0.1 \times 0.9 \cdot \theta_{99}} + \underbrace{0.1 (0.9)^2 \theta_{98}} + \underbrace{0.1 (0.9)^3 \theta_{97}} + \underbrace{0.1 (0.9)^4 \theta_{96}} + \dots \end{aligned}$$

$$\underbrace{0.9^{10}}_{\approx 0.35} \approx \frac{1}{e}$$

$$\frac{(1-\epsilon)^{1/\epsilon}}{0.9} = \frac{1}{e}$$

0.98?

$$\epsilon = 0.02 \rightarrow \underbrace{0.98^{50}}_{\approx \frac{1}{e}}$$



$$\approx \frac{1}{1-\beta}$$

$$\Sigma = 1-\beta$$

$$0.1\theta_{98} + 0.9v_{97}$$

Implementing exponentially weighted averages



$$v_0 = 0$$

$$v_1 = \beta v_0 + (1 - \beta) \theta_1$$

$$v_2 = \beta v_1 + (1 - \beta) \theta_2$$

$$v_3 = \beta v_2 + (1 - \beta) \theta_3$$

...

$$V_0 := 0$$

$$V_0 := \beta v + (1 - \beta) \theta_1$$

$$V_0 := \beta v + (1 - \beta) \theta_2$$

⋮

$$\rightarrow V_0 = 0$$

Repeat {

Get next θ_t

$$V_0 := \beta V_0 + (1 - \beta) \theta_t \leftarrow$$

}




deeplearning.ai

Optimization Algorithms

Bias correction
in exponentially
weighted average



Bias correction



$$\rightarrow v_t = \beta v_{t-1} + (1 - \beta) \theta_t$$

$$v_0 = 0$$

$$v_1 = \cancel{0.98 v_0} + 0.02 \theta_1$$

$$\begin{aligned} v_2 &= 0.98 v_1 + 0.02 \theta_2 \\ &= 0.98 \times 0.02 \times \theta_1 + 0.02 \theta_2 \\ &= 0.0196 \theta_1 + 0.02 \theta_2 \end{aligned}$$

$$\frac{v_t}{1 - \beta^t}$$

$$t=2: 1 - \beta^t = 1 - (0.98)^2 = 0.0396$$

$$\frac{v_2}{0.0396} = \frac{0.0196 \theta_1 + 0.02 \theta_2}{0.0396}$$




deeplearning.ai

Optimization Algorithms

Gradient descent with momentum



Gradient descent example



Momentum:

On iteration t :

Compute $\Delta W, \Delta b$ on current mini-batch.

☐ $V_{\Delta W} = \beta V_{\Delta W} + (1-\beta) \Delta W$

☐ $V_{\Delta b} = \beta V_{\Delta b} + (1-\beta) \Delta b$

friction \rightarrow \uparrow velocity

\uparrow acceleration

$$W := W - \alpha V_{\Delta W}, \quad b := b - \alpha V_{\Delta b}$$

$$V_{\Theta} = \beta V_{\Theta} + (1-\beta) \Theta_t$$

Implementation details

$$v_{dw} = 0, v_{db} = 0$$

On iteration t :

Compute dW, db on the current mini-batch

$$\begin{aligned} \rightarrow v_{dw} &= \beta v_{dw} + (1 - \beta) dW \\ \rightarrow v_{db} &= \beta v_{db} + (1 - \beta) db \end{aligned} \quad \left| \quad \underline{v_{dw}} = \beta v_{dw} + dW \leftarrow$$

$$\rightarrow W = W - \alpha v_{dw}, \quad b = \underline{b} - \alpha v_{db}$$

$$\frac{v_{dw}}{1 - \beta^t}$$

Hyperparameters: α, β

$$\beta = 0.9$$

average over last ≈ 10 gradients




deeplearning.ai

Optimization
Algorithms


RMSprop



0:04 / 7:41



RMSprop



On iteration t :

Compute dw, db on current mini-batch

$$\underline{S_{dw}} = \beta_2 \underline{S_{dw}} + (1 - \beta_2) \underbrace{dw^2}_{\text{element-wise}} \leftarrow \text{small}$$

$$\rightarrow \underline{S_{db}} = \beta_2 \underline{S_{db}} + (1 - \beta_2) \underline{db^2} \leftarrow \text{large}$$

$$w := w - \frac{dw}{\sqrt{\underline{S_{dw}} + \epsilon}} \leftarrow$$

$$b := b - \frac{db}{\sqrt{\underline{S_{db}} + \epsilon}} \leftarrow$$

$$\epsilon = 10^{-8}$$



deeplearning.ai

Optimization
Algorithms

Adam optimization
algorithm



Adam optimization algorithm

$$V_{dw} = 0, S_{dw} = 0, V_{db} = 0, S_{db} = 0$$

On iteration t :

Compute dw, db using current mini-batch

□ $V_{dw} = \beta_1 V_{dw} + (1 - \beta_1) dw, V_{db} = \beta_1 V_{db} + (1 - \beta_1) db \leftarrow \text{"momentum"} \beta_1$
 $S_{dw} = \beta_2 S_{dw} + (1 - \beta_2) dw^2, S_{db} = \beta_2 S_{db} + (1 - \beta_2) db^2 \leftarrow \text{"RMSprop"} \beta_2$

□ $V_{dw}^{\text{corrected}} = V_{dw} / (1 - \beta_1^t), V_{db}^{\text{corrected}} = V_{db} / (1 - \beta_1^t)$

□ $S_{dw}^{\text{corrected}} = S_{dw} / (1 - \beta_2^t), S_{db}^{\text{corrected}} = S_{db} / (1 - \beta_2^t)$

$$W := W - \alpha \frac{V_{dw}^{\text{corrected}}}{\sqrt{S_{dw}^{\text{corrected}} + \epsilon}}$$

$$b := b - \alpha \frac{V_{db}^{\text{corrected}}}{\sqrt{S_{db}^{\text{corrected}} + \epsilon}}$$

Hyperparameters choice:

☐ $\rightarrow \alpha$: needs to be tune
 $\rightarrow \beta_1 : 0.9 \rightarrow (\underline{dw})$
 $\rightarrow \beta_2 : 0.999 \rightarrow (\underline{dw^2})$
 $\rightarrow \epsilon : 10^{-8}$

Adam: Adaptive moment estimation



Adam Coates



6:46 / 7:07



Andrew Ng



deeplearning.ai


Optimization
Algorithms

Learning rate
decay



Learning rate decay

☐ Slowly reduce α




Learning rate decay

1 epoch = 1 pass through data.


□
$$\alpha = \frac{1}{1 + \text{decay-rate} * \text{epoch-num}} \alpha_0$$

□

Epoch	α
1	0.1
2	0.67
3	0.5
4	0.4
⋮	⋮

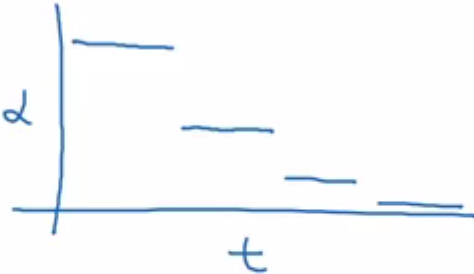


$\alpha_0 = 0.2$
 $\text{decay-rate} = 1$



Other learning rate decay methods

formula

$$\alpha = 0.95^{\text{epoch-num}} \cdot \alpha_0 \quad - \text{exponentially decay.}$$
$$\alpha = \frac{k}{\sqrt{\text{epoch-num}}} \cdot \alpha_0 \quad \text{or} \quad \frac{k}{\sqrt{t}} \cdot \alpha_0$$


discrete staircase

Manual decay.



5:35 / 6:44



Andrew Ng




deeplearning.ai


Optimization Algorithms

The problem of
local optima

Local optima in neural networks



Problem of plateaus



- Unlikely to get stuck in a bad local optima
- Plateaus can make learning slow